



# Gaussian Naive Bayes Modeling Project

Eva Azazoglu, Zoey Katzive, Alex Manioudakis, Huda Saeed

---

# What is Gaussian Naive Bayes?



## Naive Bayes Models

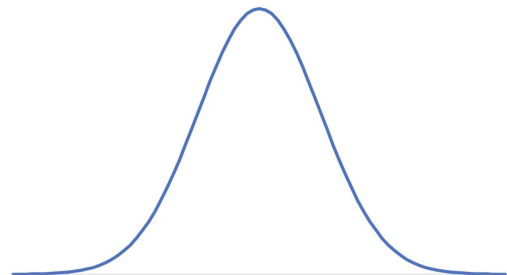
- Generative modeling: goal is to estimate the parameters of the distribution of data
- Core assumption: features are conditionally independent given the class
- Priors and likelihoods derived from Bayes' Theorem

$$P(Y|X) = \frac{P(X|Y) \cdot P(Y)}{P(X)}$$



# Gaussian Naive Bayes

- Tailored for continuous features
- Assumes features given each label follow a Gaussian (normal) distribution





## Advantages

- Computationally efficient and simple
  - Only calculates mean and variance
- Handles high-dimensional data well
  - Does not require storing or processing individual data points once parameters are learned
- Probabilistic outputs allow for nuanced interpretation of results
  - Outputs probability associated with each predicted class
- Independence assumption reduces “noise” from less relevant features



# Limitations

- Does not capture dependencies between features
  - Conditional independence assumption does not hold in real-life data
- Performance weakens when distribution is not perfectly normal or decision boundaries between classes are complex
- Sensitive to outliers
- May not identify minority classes accurately when the classes are imbalanced
  - Prior probability  $P(Y)$  influences the posterior probability

---

# The Math Behind the Algorithm



# Representation

Consists of the domain, range, sampling method, and function

1. All  $d$  features assumed to be normally distributed and independent given the class label  $\rightarrow X \in \mathbb{R}^d$
2. The range is a class label; for  $K$  labels we have  $\rightarrow Y \in \{0, 1, \dots, K\}$
3. The sample is taken from the features (normal distribution assumed conditioned on label) paired with the label outcome (unknown distribution)
4. Observation  $\rightarrow$  Output (class that has the maximum probability for an observation)





## GNB Function

- Prior probability calculated
  - Number of times that a class occurred/number of rows in the dataset
- For each feature, subset based on the class → calculate mean and standard deviation
  - Result: K\*D mean and variance pairs (K: number of classes, D: number of features)
- Observation's value for each attribute and each mean-variance pair are plugged into the Gaussian probability function:

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

$x$ : observation's value for attribute  $d$   
 $\mu$ : mean of attribute  $d$  given class  $k$   
 $\sigma^2$ : variance of attribute  $d$  given class  $k$

## Representation: Bayes Theorem

$$P(Y|X) = \frac{P(X|Y) \cdot P(Y)}{P(X)}$$

*Bayes' Theorem*



$$P(Y|x_1, x_2, \dots, x_D) \propto P(Y) \prod_{i=1}^D P(x_i|Y)$$

*Bayes' Theorem in  
our model*

$$\log P(Y) + \sum_{i=1}^D \log P(x_i|Y)$$

*Convert to log space*



$$P(Y|x_1, x_2, \dots, x_D) \propto \exp\left(\log P(Y) + \sum_{i=1}^D \log P(x_i|Y)\right)$$

*Exponentiate to convert  
to probability*



## Loss

We can define the log loss of a Gaussian Naive Bayes algorithm over all N data points:

$$\mathcal{L}(\theta) = -\frac{1}{N} \sum_{i=1}^N \log P(y_i | x_i, \theta)$$

- Theta  $\rightarrow$  the parameters of the function for Gaussian likelihood (mean and variance)
- P  $\rightarrow$  the Gaussian probability



# Optimizer

- Empirical risk minimization → want to find the minimum of the loss
- For  $m$  data points, we want

$$\arg \min_{\theta} \mathcal{L}(\theta)$$

- Note: Naive Bayes is a generative algorithm, so the loss and optimizer exist in relation to distributions of the continuous features.



# Pseudocode

GIVEN:

- `n_classes`: number of classes
- `X`: 2D array representing training features
- `y`: 1D array representing training labels
- `inputs`: 2D array representing new examples to predict



# Pseudocode

```
TRAIN(X, y):  
    Initialize priors as a 1D array of length n_classes to all zeros  
    Initialize means as a 2D array of zeros with dimensions (n_classes, d)  
    Initialize stds as a 2D array of zeros with dimensions (n_classes, d)  
  
    For each class j:  
        class_count = number of elements in y that are equal to j  
        priors[j] = class_count / (total number of elements in y)  
  
    For each class j:  
        For each feature i:  
            Determine ind as the indices of all rows in y where y == j  
            means[j, i] = the average value of feature i across ind  
            stds[j, i] = the standard deviation of feature i across ind  
  
    Return priors, means, stds
```



# Pseudocode

```
PREDICT(inputs):  
    log_priors = natural log of priors  
    Initialize predictions as empty array  
  
    For each input i:  
        Set prob = copy of log_priors  
        For each class j:  
            Add the log of the Gaussian PDF for each feature f to prob[j]  
  
            
$$-\frac{1}{2} \ln(2\pi\sigma_{j,f}^2) - \frac{(x_{i,f} - \mu_{j,f})^2}{2\sigma_{j,f}^2}$$
  
  
            prob[j] = exponential_function(prob[j])  
  
        Identify the class that corresponds to highest probability  
        Append that class to predictions  
  
    Return predictions as an array
```

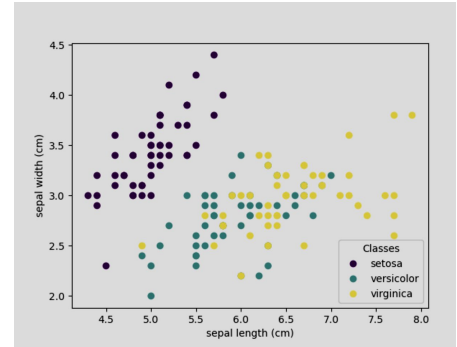
---

# Our Project



# The Iris Dataset

- Classic machine learning experimentation dataset
- Consists of 150 observations from three distinct species of iris
- Includes measurements of four key features: sepal length, sepal width, petal length, and petal width
- Each sample is accompanied by a label that indicates its species





## Reproducing a GNB Model on Iris

- Bernd Klein's "Machine Learning with Python Tutorial" textbook implements a Gaussian Naive Bayes classifier on the Iris dataset using scikit-learn
- Outputs to match: precision, recall, f1 score, support, confusion matrix
- Work to reproduce: classification report and confusion matrix

```
GaussianNB()
precision    recall  f1-score   support

0           1.00      1.00      1.00         50
1           0.94      0.94      0.94         50
2           0.94      0.94      0.94         50

avg / total         0.96      0.96      0.96        150

[[50  0  0]
 [ 0 47  3]
 [ 0  3 47]]
```



## Our Process

- Create a class GNB() with methods for train, predict, classification report, and confusion matrix
- Create unit tests to validate train and predict methods
- Run GNB() on the Iris dataset
- Verify that output matches reproduced work (it does!)

---

# Gaussian Naive Bayes in the Real World



## Other Works

“Human Activity Recognition Using Gaussian Naïve Bayes Algorithm in Smart Home” (Shen & Fang, 2020)

<https://iopscience.iop.org/article/10.1088/1742-6596/1631/1/012059>

- **IoT sensors to predict one's activities in the household**
- Used GNB to generate distributions on attributes created from sensor data
- Assumption that continuous attributes are normally distributed and independent
- Features were constructed from collected sensor data to show things like “time spent on activity” or “number of ‘on’ sensors”
- Optimal feature values are determined throughout the experiment
- 7% higher accuracy with GNB (89.5%) than with NB (82.7%)
- Impact: aging populations = increased need to cater to the at-home needs of the elderly
  - IoT can help, but ethical issues with privacy implications

---

# Takeaways



# Summary

Challenges:

- Finding work to reproduce with purely quantitative features
- Defining the loss and the optimizer for a generative model

What we found interesting:

- How viable GNB is in applications despite strong assumptions of normality + independence
- Logistic regression is often a much better alternative
- Works for discrete quantitative features

---

# Questions?